

Sharan Giri

giri.sha@northeastern.edu | (857) 565-4042 | github.com/sharan0276 | linkedin.com/in/sharan-giri | sharan0276.github.io

PROFESSIONAL STATEMENT

AI/ML Engineer with 2 years shipping production pipelines at Barclays under SLA and compliance controls. I build agentic orchestration systems where LLM outputs stay policy-bound, evaluated, and traced to source, delivering policy-gated guardrails and multi-step reasoning pipelines at production scale.

EDUCATION

Khoury College of Computer Science, Northeastern University M.S. in Data Science (GPA: 4.0) Secured Third Place at Google-Northeastern Expo for AudioSEEK (Dec 2025)	Boston, MA Sep 2024 – May 2026
Sri Sivasubramaniya Nadar College of Engineering, Anna University B.E. in Computer Science Engineering (GPA: 3.7)	India Aug 2018 – Jun 2022

EXPERIENCE

Barclays Global Service Centre – Markets And Private Banking Technology Automation Developer Data Pipeline and API Integration - Owned 3-month delivery cycles from requirements to deployment, running failure-reproduction UAT with trading teams and resolving post-launch issues within 2-day SLA - Built schema-validated SWIFT pipelines enforcing data quality, lifting throughput 50% and saving 85 hrs/month - Rebuilt 4 trader workflows as API-integrated automation after identifying unstable UI mid-project, cutting turnaround 33% - Traced ML extraction failures to upstream format inconsistency, coordinated with trading teams to enforce a structured reply standard at source - Implemented queue-based HITL escalation across 6 regulated workflows, routing compliance-threshold exceptions to ops desks on sign-off	Aug 2022 – Jun 2024 India
---	------------------------------

PROJECTS

CourseLENS Governed LLM Inference System with Policy Controls - Parsed PPTX, PDF, and web sources into a normalized schema with skip-and-log failure isolation per source - Designed two-phase ingestion decoupling extraction from embedding with fault-tolerant JSON checkpoint restart - Engineered a 4-stage agentic state machine enforcing answer-release policy through staged reasoning, blocking 15/16 adversarial bypass attempts - Designed dual-model architecture with an independent LLM assessor calibrating state transitions from user inputs - Held 96.15% faithfulness via Ragas, logging every agent decision and state transition to a full audit trail	Jan 2026 – Apr 2026
FinAgent Grounded Comparative Risk Inference Platform - Built a two-stage pipeline gating LLM inference behind deterministic FinBERT risk classification, cutting token costs 98% - Aligned 5-year risk trajectories across 72 SEC filings and 16 comparable companies for year-over-year peer analysis - Stored classified risk signals by company and fiscal year in a vector store, enabling time-aware peer matching on query - Traced every risk signal to its source SEC filing via deterministic chunk IDs, enabling full audit trail on analyst queries - Served FAISS retrieval over BGE-M3 embeddings through FastAPI endpoints and a Streamlit analyst interface	Mar 2026 – Apr 2026
AudioSEEK Unstructured Audio Intelligence Platform <i>Google-Northeastern Expo</i> - Built a retrieval platform over 30+ hours of unstructured audio, returning timestamp-grounded answers in sub-15s - Selected best transcription across 3 models via WER and ROUGE-L, enforcing accuracy-based ingestion decisions - Enforced temporal constraints on a hybrid SQL-FAISS store, anchoring every retrieved chunk to its source timestamp - Operationalized token usage, latency, and drift tracking via MLflow on GKE for production agent monitoring	Sep 2025 – Dec 2025
Demand Signal Modeling Explainable AI Ranking Engine over Behavioral Signals - Built an XGBoost ranking engine over a structured JSON review dataset, achieving 89% Precision@100 on high-potential candidate identification - Extracted behavioral signals across momentum trajectories, review velocity, and user sentiment using VADER NLP - Ran SHAP audits to identify and purge proxy leakage, isolating momentum and sentiment as the dominant predictive drivers - Validated global ranking quality via NDCG (0.9596), proving ranking intelligence beyond pointwise classification	Nov 2025 – Dec 2025

TECHNICAL SKILLS

AI Systems & Retrieval: RAG Pipelines, Agentic AI Systems, LangChain, LangGraph, FAISS, Vector Databases, BM25, BGE-M3, HuggingFace, Pydantic, Structured Output Generation, Prompt Engineering
Evaluation & Observability: Ragas, MLflow, WER, ROUGE-L, LLM Evaluation, LLM Optimization, ML Observability, SHAP
Data & Infrastructure: Python, SQL, PySpark, Pandas, Airflow, FastAPI, Streamlit, REST APIs, DuckDB, Git, CI/CD (GitHub Actions), Docker, Kubernetes (GKE), Terraform, GCP (Cloud Run, GCS, Firestore)